

## 16.4.2

## Spark DataFrames and Datasets

**It's** time to start working with Spark! You'll start by using DataFrames. You will soon see that Spark has a lot of similarities to the Pandas library, which you have used before.

Working in Spark requires us to put data into DataFrames. If you're wondering if these DataFrames are comparable to those in Pandas, you're correct—Spark DataFrames are very similar. Just as in Pandas, the first step is to load your data into a DataFrame.

 [Retake](#)

The schema, or structure, for DataFrames and datasets contains the column names and the data types contained within.

The schema can be inferred automatically by letting Spark determine the schema on its own when data is read in or you can define the schema manually and tell spark to use that.

In this module, we'll use the Python version of Spark, PySpark.

 [Retake](#)

Using your Google Colab Notebook, with PySpark installed, follow along with the code.

Create a Spark session by importing the library and setting the `spark` variable to the code below:

```
# Start Spark session
from pyspark.sql import SparkSession
spark = SparkSession.builder.appName("DataFrameBasics").getOrCreate()
```

## NOTE

Every time you open a new Google Colab Notebook, run the Spark installation and then start a Spark session. The app name can be different for each notebook.

This creates a Spark application called "DataFrameBasics."

Spark enables us to create a DataFrame from scratch by passing in a list of tuples to the `createDataFrame` method followed by a list of the column names. The `show` method will display the DataFrame, which is similar to using the `head()` function in Pandas. Enter and run the following code:

```
dataframe = spark.createDataFrame([
    (0, "Here is our DataFrame"),
    (1, "We are making one from scratch"),
    (2, "This will look very similar to a Pandas DataFrame")
], ["id", "words"])

dataframe.show()
```

```
+---+-----+
| id |      words |
|  0 | Here is out DataF... |
|  1 | We are making on... |
|  2 | This will look ve... |
+---+-----+
```

Spark also lets us import data directly into a DataFrame. To do this, we import `SparkFiles` from the pyspark library that allows us to retrieve files. The next three lines of code tell Spark to pull data from Amazon's Simple Storage Service (S3), a cloud-based data storage service. This boilerplate code can be used to read other public files hosted on Amazon's services. We'll dive more into cloud storage later on in the module.

In the next cell in your Colab notebook, run the following code:

```
# Read in data from S3 Buckets
from pyspark import SparkFiles
url = "https://s3.amazonaws.com/dataviz-curriculum/day_1/food.csv"
spark.sparkContext.addFile(url)
df = spark.read.csv(SparkFiles.get("food.csv"), sep=",", header=True)
```

Type the following code to use `show()` again to display the results, as follows:

```
# Show DataFrame
df.show()
```

food	price
pizza	0
sushi	12
chinese	10

Spark will infer the schema from the data, unless otherwise specified. We can check the schema by running the following code:

```
# Print our schema
df.printSchema()
```

Spark also allows users to view columns and a dataset description by running each of the code blocks:

```
# Show the columns
df.columns
```

```
['food', 'price']
```

```
# Describe our data
df.describe()
```

```
DataFrame[summary: string, food: string, price: string]
```

Notice that the DataFrame is claiming that `price` is a string. Generally, `price` is either stored as an integer or floating-point number, so you'll need to change this column.

In this case, we can set our schema and then apply it to the data. We'll start by importing the different types of data with the following code:

```
# Import struct fields that we can use
from pyspark.sql.types import StructField, StringType, IntegerType, StructType
```

Next, create the schema by creating a `StructType`, which is one of Spark's complex types, like an array or map. The `StructField` will define the column name, the data type held, and a Boolean to define whether null values will be included or not:

```
# Next we need to create the list of struct fields
schema = [StructField("food", StringType(), True), StructField("price", IntegerType(), True),]
schema
```

Next, enter the code that will pass the schema just created as fields in a `StructType`. All this will be stored in a variable called `final`:

```
# Pass in our fields
final = StructType(fields=schema)
final
```

Now that we have a predefined schema, we can read in the data again, only this time passing in our own schema. Type and run the following code in a new notebook cell:

```
# Read our data with our new schema
dataframe = spark.read.csv(SparkFiles.get("food.csv"), schema=final, sep=",", header=True)
dataframe.printSchema()
```

There are a few different ways to access our data with Spark. Run the following commands and look at the results:

```
dataframe['price']  
  
Column<b 'price'>  
  
type(dataframe['price'])  
  
pyspark.sql.column.Column  
  
dataframe.select('price')  
  
DataFrame[price: int]  
  
type(dataframe.select('price'))  
  
pyspark.sql.dataframe.DataFrame  
  
dataframe.select('price').show()  
  
+-----+  
|       |  
+-----+  
| price |  
|      0 |  
|     12 |  
|     10 |  
+-----+
```

Again, you may notice some similarities to Pandas. For example, in both Pandas and Spark, you can select a column using the DataFrame name, followed by the column's name in square brackets. In Pandas, you can quickly take a look at a DataFrame using `head()`; in Spark, you can do something similar using `show()`.

#### NOTE

You might notice that code like `dataframe['price']` isn't performing as expected. After running this code, we get the column name, but no results until the `show()` function runs. `show` is an action, whereas `select` is a transformation. We'll cover what this means in the next section.

We can manipulate columns in Spark as well. Run the code below and guess what will be displayed:

```
# Add new column  
dataframe.withColumn('newprice', dataframe['price']).show()  
  
# Update column name  
dataframe.withColumnRenamed('price', 'newerprice').show()
```

```
# Double the price
dataframe.withColumn('doubleprice',dataframe['price']*2).show()

# Add a dollar to the price
dataframe.withColumn('add_one_dollar',dataframe['price']+1).show()

# Half the price
dataframe.withColumn('half_price',dataframe['price']/2).show()
```

The first cell duplicates the `price` column into a new column, preserving all its rows, and naming that column `newprice`. The second cell simply renames the `price` column as `newprice`.

```
# Add new column
dataframe.withColumn('newprice', dataframe['price']).show()
```

food	price	newprice
pizza	0	0
sushi	12	12
chinese	10	10

```
# Update column name
dataframe.withColumnRenamed('price', 'newerprice').show()
```

food	newerprice
pizza	0
sushi	12
chinese	10

The next three cells created a new column, but they performed some operation on the original:

```
# Double the price
dataframe.withColumn('doubleprice', dataframe['price']*2).show()
```

food	price	doubleprice
pizza	0	0
sushi	12	24
chinese	10	20

```
# Add a dollar to the price
dataframe.withColumn('add_one_dollar', dataframe['price']+1).show()
```

food	price	add_one_dollar
pizza	0	1
sushi	12	13
chinese	10	11

```
# Half the price
dataframe.withColumn('half_price', dataframe['price']/2).show()
```

food	price	half_price
pizza	0	0.0
sushi	12	6.0
chinese	10	5.0

Next, we'll dive into Spark functions and why some of these commands produced actual results while others just relayed information about our data.

© 2020 - 2022 Trilogy Education Services, a 2U, Inc. brand. All Rights Reserved.