

16.9.1

PySpark ETL

You and Jennifer are excited about all the new technologies you're ready to use. You explain to Jennifer that you can now move your ETL process to the cloud to take the load off your local laptops, use strictly cloud resources, and eventually lead to automation. Let's get started with extraction!

Let's run through a mock scenario using two different types of raw data stored in S3. Our goal is to get this raw data from S3 into an RDS database. Let's start with uploading files to your own S3 bucket:

Create an S3 bucket, and then load the following files into the bucket:

- [user_data.csv](https://2u-data-curriculum-team.s3.amazonaws.com/dataviz-online/module_16/user_data.csv) [_](https://2u-data-curriculum-team.s3.amazonaws.com/dataviz-online/module_16/user_data.csv)(https://2u-data-curriculum-team.s3.amazonaws.com/dataviz-online/module_16/user_data.csv)
- [user_payment.csv](https://2u-data-curriculum-team.s3.amazonaws.com/dataviz-online/module_16/user_payment.csv) [_](https://2u-data-curriculum-team.s3.amazonaws.com/dataviz-online/module_16/user_payment.csv)(https://2u-data-curriculum-team.s3.amazonaws.com/dataviz-online/module_16/user_payment.csv)

IMPORTANT

Remember to make the bucket and files public.

Assume your company already has three tables set up in the RDS database and would like to get the raw data from S3 into the database. Create a new database in pgAdmin called "my_data_class_db." We'll have it represent the company database by first running the following schema in pgAdmin for our RDS:

```
-- Create Active User Table
CREATE TABLE active_user (
  id INT PRIMARY KEY NOT NULL,
  first_name TEXT,
  last_name TEXT,
  username TEXT
);
```

```
CREATE TABLE billing_info (  
  billing_id INT PRIMARY KEY NOT NULL,  
  street_address TEXT,  
  state TEXT,  
  username TEXT  
);  
  
CREATE TABLE payment_info (  
  billing_id INT PRIMARY KEY NOT NULL,  
  cc_encrypted TEXT  
);
```

NOTE

Table creation is not part of the ETL process. We're creating the tables to represent a pre-established database you need for the raw data. In a real-life situation, databases will already have a well-defined schema and tables for you, as the engineer, to process data into.

Start with creating a new notebook, installing Spark:

```
import os  
# Find the latest version of spark 3.0 from http://www.apache.org/dist/spark/ and enter as the spark version  
# For example:  
# spark_version = 'spark-3.0.3'  
spark_version = 'spark-3.<enter version>'  
os.environ['SPARK_VERSION']=spark_version  
  
# Install Spark and Java  
!apt-get update  
!apt-get install openjdk-11-jdk-headless -qq > /dev/null  
!wget -q http://www.apache.org/dist/spark/$SPARK_VERSION/$SPARK_VERSION-bin-hadoop2.7.tgz  
!tar xf $SPARK_VERSION-bin-hadoop2.7.tgz  
!pip install -q findspark  
  
# Set Environment Variables  
import os  
os.environ["JAVA_HOME"] = "/usr/lib/jvm/java-11-openjdk-amd64"  
os.environ["SPARK_HOME"] = f"/content/{spark_version}-bin-hadoop2.7"  
  
# Start a SparkSession
```

```
import findspark
findspark.init()
```

We'll use Spark to write directly to our Postgres database. But in order to do so, there are few more lines of code we need.

First, enter the following code to download a Postgres driver that will allow Spark to interact with Postgres:

```
!wget https://jdbc.postgresql.org/download/postgresql-42.2.16.jar
```

You should get a message containing the words "HTTP request sent, awaiting response... 200 OK," indicating that your request was processed without a problem.

Then, start a Spark session with an additional option that adds the driver to Spark:

```
from pyspark.sql import SparkSession
spark = SparkSession.builder.appName("CloudETL").config("spark.driver.extraClassPath", "/content/postgresql-42
```

We have performed the first two steps of the ETL process before with PySpark, so let's quickly review those.

Extract

We can connect to data storage, then extract that data into a DataFrame. We'll do this on two datasets, and be sure to replace the bucket name with one of your own.

We'll start by importing SparkFiles from PySpark into our notebook. This will allow Spark to add a file to our Spark project.

Next, the file is read in with the `read` method and combined with the `csv()` method, which pulls in our CSV stored in SparkFiles and infers the schema. `SparkFiles.get()` will have Spark retrieve the specified file, since we are dealing with a CSV. The `","` is the chosen separator, and we will have Spark determine the head for us. Enter the following code:

```
# Read in data from S3 Buckets
from pyspark import SparkFiles
```

```
url = "https://YOUR-BUCKET-NAME.s3.amazonaws.com/user_data.csv"
spark.sparkContext.addFile(url)
user_data_df = spark.read.csv(SparkFiles.get("user_data.csv"), sep=",", header=True, inferSchema=True)
```

Finally, an action is called to show the first 10 runs and confirm our data extraction by entering the following code:

```
# Show DataFrame
user_data_df.show()
```

Repeat a similar process to load in the other data. Enter the code:

```
url = "https://YOUR-BUCKET-NAME.s3.amazonaws.com/user_payment.csv"
spark.sparkContext.addFile(url)
user_payment_df = spark.read.csv(SparkFiles.get("user_payment.csv"), sep=",", header=True, inferSchema=True)

# Show DataFrame
user_payment_df.show()
```

Transform

Now that the raw data stored in S3 is available in a PySpark DataFrame, we can perform our transformations.

First, join the two tables:

```
# Join the two DataFrame
joined_df= user_data_df.join(user_payment_df, on="username", how="inner")
joined_df.show()
```

username	id	first_name	last_name	active_user	street_address	state	billing_id	cc_encrypted
ibearham0	1	Cletus	Lithcow	FALSE	78309 Riverside Way	Virginia	1	a799fcafe47d7fb19...
wwaller1	2	Caz	Felgat	FALSE	83 Hazelcrest Place	Alabama	2	a799fcafe47d7fb19...
ichesn2	3	Kerri	Crowson	FALSE	112 Eliot Pass	North Carolina	3	a799fcafe47d7fb19...
tsnarr3	4	Freddie	Caghy	FALSE	15 Merchant Way	New York	4	a799fcafe47d7fb19...
fwherrit4	5	Sadella	Deuss	FALSE	079 Acker Avenue	Tennessee	5	a799fcafe47d7fb19...
fstappard5	6	Fraser	Korneev	TRUE	76084 Novick Court	Minnesota	6	a799fcafe47d7fb19...
lhambling6	7	Demott	Rapson	TRUE	86320 Dahle Park	District of Columbia	7	a799fcafe47d7fb19...
drude7	8	Robert	Poile	FALSE	1540 Manitowish Hill	Georgia	8	a799fcafe47d7fb19...
bspawton8	9	Nollie	null	TRUE	4 Katie Court	Ohio	9	a799fcafe47d7fb19...
rmackeller9	10	Merilyn	Frascone	FALSE	387 Duke Street	Ohio	10	a799fcafe47d7fb19...
cdennerleya	11	Rickie	Tredwell	FALSE	04 Monterey Center	Missouri	11	a799fcafe47d7fb19...
gsarfasb	12	Charmane	Connerry	FALSE	0 Larry Junction	Florida	12	a799fcafe47d7fb19...
mpichefordc	13	Nerti	Kerins	FALSE	68 Portage Trail	California	13	a799fcafe47d7fb19...
bingryd	14	Bart	null	FALSE	8 Homewood Court	District of Columbia	14	a799fcafe47d7fb19...
whainerte	15	Sadella	Jaram	TRUE	7528 Waxwing Terrace	Connecticut	15	a799fcafe47d7fb19...
mdrewetf	16	Dicky	Runnett	FALSE	1793 Delaware Park	Florida	16	a799fcafe47d7fb19...
droughsedg	17	Hewitt	Trammel	TRUE	2455 Corry Alley	North Carolina	17	a799fcafe47d7fb19...
abaakeh	18	Gilligan	Boys	FALSE	2 Raven Court	Florida	18	a799fcafe47d7fb19...
ydudeniei	19	Ted	Knowlys	TRUE	31 South Drive	Ohio	19	a799fcafe47d7fb19...
ckermittj	20	Darb	Carrel	FALSE	406 Park Meadow C...	Minnesota	20	a799fcafe47d7fb19...

Next, drop any rows with null or "not a number" (NaN) values:

```
# Drop null values
dropna_df = joined_df.dropna()
dropna_df.show()
```

username	id	first_name	last_name	active_user	street_address	state	billing_id	cc_encrypted
ibearham0	1	Cletus	Lithcow	FALSE	78309 Riverside Way	Virginia	1	a799fcafe47d7fb19...
wwaller1	2	Caz	Felgat	FALSE	83 Hazelcrest Place	Alabama	2	a799fcafe47d7fb19...
ichesn2	3	Kerri	Crowson	FALSE	112 Eliot Pass	North Carolina	3	a799fcafe47d7fb19...
tsnarr3	4	Freddie	Caghy	FALSE	15 Merchant Way	New York	4	a799fcafe47d7fb19...
fwherrit4	5	Sadella	Deuss	FALSE	079 Acker Avenue	Tennessee	5	a799fcafe47d7fb19...
fstappard5	6	Fraser	Korneev	TRUE	76084 Novick Court	Minnesota	6	a799fcafe47d7fb19...
lhambling6	7	Demott	Rapson	TRUE	86320 Dahle Park	District of Columbia	7	a799fcafe47d7fb19...
drude7	8	Robert	Poile	FALSE	1540 Manitowish Hill	Georgia	8	a799fcafe47d7fb19...
rmackeller9	10	Merilyn	Frascone	FALSE	387 Duke Street	Ohio	10	a799fcafe47d7fb19...
cdennerleya	11	Rickie	Tredwell	FALSE	04 Monterey Center	Missouri	11	a799fcafe47d7fb19...
gsarfasb	12	Charmane	Connerry	FALSE	0 Larry Junction	Florida	12	a799fcafe47d7fb19...
mpichefordc	13	Nerti	Kerins	FALSE	68 Portage Trail	California	13	a799fcafe47d7fb19...
whainerte	15	Sadella	Jaram	TRUE	7528 Waxwing Terrace	Connecticut	15	a799fcafe47d7fb19...
mdrewetf	16	Dicky	Runnett	FALSE	1793 Delaware Park	Florida	16	a799fcafe47d7fb19...
droughsedg	17	Hewitt	Trammel	TRUE	2455 Corry Alley	North Carolina	17	a799fcafe47d7fb19...
abaakeh	18	Gilligan	Boys	FALSE	2 Raven Court	Florida	18	a799fcafe47d7fb19...
ydudeniei	19	Ted	Knowlys	TRUE	31 South Drive	Ohio	19	a799fcafe47d7fb19...
ckermittj	20	Darb	Carrel	FALSE	406 Park Meadow C...	Minnesota	20	a799fcafe47d7fb19...
ipowisk	21	Diandra	Cancellor	FALSE	1 Fisk Parkway	North Carolina	21	a799fcafe47d7fb19...
dtalton1	22	Ulrika	Itzhayek	FALSE	890 Lakewood Alley	California	22	a799fcafe47d7fb19...

Filter for active users:

```
# Load in a sql function to use columns
from pyspark.sql.functions import col

# Filter for only columns with active users
cleaned_df = dropna_df.filter(col("active_user") == True)
cleaned_df.show()
```

username	id	first_name	last_name	active_user	street_address	state	billing_id	cc_encrypted
fstappard5	6	Fraser	Korneev	TRUE	76084 Novick Court	Minnesota	6	a799fcafe47d7fb19...
lhambling6	7	Demott	Rapson	TRUE	86320 Dahle Park	District of Columbia	7	a799fcafe47d7fb19...
wheinerte	15	Sadella	Jaram	TRUE	7528 Waxwing Terrace	Connecticut	15	a799fcafe47d7fb19...
droughsedgeg	17	Hewitt	Trammel	TRUE	2455 Corry Alley	North Carolina	17	a799fcafe47d7fb19...
ydudeniei	19	Ted	Knowlys	TRUE	31 South Drive	Ohio	19	a799fcafe47d7fb19...
fmyttonm	23	Annmarie	Lafond	TRUE	35 Oriole Place	Georgia	23	a799fcafe47d7fb19...
bfletcher	28	Toma	Sokell	TRUE	39641 Eggendart Hill	Maryland	28	a799fcafe47d7fb19...
gturleyt	30	Ram	Lefever	TRUE	9969 Laurel Alley	Texas	30	a799fcafe47d7fb19...
calyukinu	31	Raddie	Heindle	TRUE	811 Talmadge Road	Ohio	31	a799fcafe47d7fb19...
ckleinlererw	33	Wallie	Caws	TRUE	9999 Kenwood Pass	Oregon	33	a799fcafe47d7fb19...
pshanklandx	34	Derril	Varfolomeev	TRUE	4 Jenifer Court	Florida	34	a799fcafe47d7fb19...
enelanel2	39	Kelcy	Wheway	TRUE	93207 Morningstar...	Florida	39	a799fcafe47d7fb19...
sfollet13	40	Dorree	Rookeby	TRUE	2 Troy Circle	California	40	a799fcafe47d7fb19...
mtesh14	41	Martyn	Tott	TRUE	728 Muir Lane	Florida	41	a799fcafe47d7fb19...
tseyfart16	43	Cally	Thody	TRUE	1 Graceland Plaza	Florida	43	a799fcafe47d7fb19...
hfarrier18	45	Ted	Pittaway	TRUE	767 Little Fleur ...	North Carolina	45	a799fcafe47d7fb19...
nabbie1b	48	Fifi	Lidgley	TRUE	6744 Sutherland Road	South Carolina	48	a799fcafe47d7fb19...
ystadding1d	50	Ashely	O'Hern	TRUE	929 Scoville Park	Florida	50	a799fcafe47d7fb19...
hhallgalley1g	53	Diannne	Osbaldeston	TRUE	0 Mesta Pass	Tennessee	53	a799fcafe47d7fb19...
ageaveny1n	60	Sonny	Jeskin	TRUE	50 Sutherland Drive	Massachusetts	60	a799fcafe47d7fb19...

Next, select columns to create three different DataFrames that match what is in the AWS RDS database. Create a DataFrame to match the active_user table:

```
# Create user dataframe to match active_user table
clean_user_df = cleaned_df.select(["id", "first_name", "last_name", "username"])
clean_user_df.show()
```

id	first_name	last_name	username
6	Fraser	Korneev	fstappard5
7	Demott	Rapson	lhambling6
15	Sadella	Jaram	wheinerte
17	Hewitt	Trammel	droughsedgeg
19	Ted	Knowlys	ydudeniei
23	Annmarie	Lafond	fmyttonm
28	Toma	Sokell	bfletcher
30	Ram	Lefever	gturleyt
31	Raddie	Heindle	calyukinu
33	Wallie	Caws	ckleinlererw
34	Derril	Varfolomeev	pshanklandx
39	Kelcy	Wheway	enelanel2
40	Dorree	Rookeby	sfollet13
41	Martyn	Tott	mtesh14
43	Cally	Thody	tseyfart16
45	Ted	Pittaway	hfarrier18
48	Fifi	Lidgley	nabbie1b
50	Ashely	O'Hern	ystadding1d
53	Diannne	Osbaldeston	hhallgalley1g
60	Sonny	Jeskin	ageaveny1n

Next, create a DataFrame to match the billing_info table:

```
# Create user dataframe to match billing_info table
clean_billing_df = cleaned_df.select(["billing_id", "street_address", "state", "username"])
clean_billing_df.show()
```

billing_id	street_address	state	username
6	76084 Novick Court	Minnesota	fstappard5
7	86320 Dahle Park	District of Columbia	lhambling6
15	7528 Waxwing Terrace	Connecticut	wheinerte
17	2455 Corry Alley	North Carolina	droughsedgex
19	31 South Drive	Ohio	ydudeniei
23	35 Oriole Place	Georgia	fmyttonm
28	39641 Eggendart Hill	Maryland	bfletcherr
30	9969 Laurel Alley	Texas	gturleyt
31	811 Talmadge Road	Ohio	calyukinu
33	9999 Kenwood Pass	Oregon	ckleinlererw
34	4 Jenifer Court	Florida	pshanklandx
39	93207 Morningstar...	Florida	enelanel2
40	2 Troy Circle	California	sfollet13
41	728 Muir Lane	Florida	mtesh14
43	1 Graceland Plaza	Florida	tseyfart16
45	767 Little Fleur ...	North Carolina	hfarrrier18
48	6744 Sutherland Road	South Carolina	nabbie1b
50	929 Scoville Park	Florida	ystadding1d
53	0 Mesta Pass	Tennessee	hhallgalley1g
60	50 Sutherland Drive	Massachusetts	ageaveny1n

Finally, create a DataFrame to match the payment_info table:

```
# Create user dataframe to match payment_info table
clean_payment_df = cleaned_df.select(["billing_id", "cc_encrypted"])
clean_payment_df.show()
```

billing_id	cc_encrypted
6	a799fcafe47d7fb19...
7	a799fcafe47d7fb19...
15	a799fcafe47d7fb19...
17	a799fcafe47d7fb19...
19	a799fcafe47d7fb19...
23	a799fcafe47d7fb19...
28	a799fcafe47d7fb19...
30	a799fcafe47d7fb19...
31	a799fcafe47d7fb19...
33	a799fcafe47d7fb19...
34	a799fcafe47d7fb19...
39	a799fcafe47d7fb19...
40	a799fcafe47d7fb19...
41	a799fcafe47d7fb19...
43	a799fcafe47d7fb19...
45	a799fcafe47d7fb19...
48	a799fcafe47d7fb19...
50	a799fcafe47d7fb19...
53	a799fcafe47d7fb19...
60	a799fcafe47d7fb19...

Once our data has been transformed to fit the tables in our database, we're ready to move on to the "Load" step.

Load

The final step is to get our transformed raw data into our database. PySpark can easily connect to a database to load the DataFrames into the table. First, we'll do some configuration to allow the connection with the following code:

```
# Store environmental variable
from getpass import getpass
password = getpass('Enter database password')
# Configure settings for RDS
mode = "append"
jdbc_url="jdbc:postgresql://<connection string>:5432/<database-name>"
config = {"user": "postgres",
          "password": password,
          "driver": "org.postgresql.Driver"}
```

You'll need to provide your username and password, and also supply the AWS server name where `<connection string>` is located in the code above. To find it in PgAdmin, right-click AWS in the Server directory listing on the left side of PgAdmin, and then select Properties in the drop-down menu. Select the Connection tab in the window that opens, and then select the address in the Host name/address field. Copy that address and paste it in place of `<connection string>`.

Let's further break down what's happening here:

- The `getpass` module is used to temporarily store your database password without exposing it to the public.
- `mode` is what we want to do with the DataFrame to the table, such as `overwrite` or `append`. We'll append to the current table because every time we run this ETL process, we'll want more data added to our database without removing any.
- The `jdbc_url` is the connection string to our database.
 - Replace `<connection string>` with the endpoint connection url found from your AWS RDS console.
 - Replace `<database name>` with the name of your database you wish to connect to.
- A dictionary of configuration that includes the `user`, `password`, and `driver` to what type of database is being used.
 - The `user` field is the username for your database, which should be `postgres` if you followed with the creation of the RDS instance. Otherwise, enter the one you created.
 - The `password` would be the password you created when making the RDS instance.

NOTE

If you forget anything like the name of the database or user name you can check on pgAdmin for these values. Be sure that you are entering the name of the database and not the name of your server in the connection string.

The cleaned DataFrames can then be written directly to our database by using the `.write.jdbc` method that takes in the parameters we set:

- The connection string stored in `jdbc_url` is passed to the URL argument.
- The corresponding name of the table we are writing the DataFrame to.
- The mode we're using, which is "append."
- The connection configuration we set up passed to the properties.

The code is as follows:

```
# Write DataFrame to active_user table in RDS
clean_user_df.write.jdbc(url=jdbc_url, table='active_user', mode=mode, properties=config)
```

```
# Write dataframe to billing_info table in RDS
clean_billing_df.write.jdbc(url=jdbc_url, table='billing_info', mode=mode, properties=config)
```

```
# Write dataframe to payment_info table in RDS
clean_payment_df.write.jdbc(url=jdbc_url, table='payment_info', mode=mode, properties=config)
```

Let's wrap up by double-checking our work and running queries in pgAdmin on our database to confirm that the load did exactly what we wanted:

```
-- Query database to check successful upload
SELECT * FROM active_user;
SELECT * FROM billing_info;
SELECT * FROM payment_info;
```

Nice work! You now have enough knowledge and practice with PySpark and AWS to begin your client project.

© 2020 - 2022 Trilogy Education Services, a 2U, Inc. brand. All Rights Reserved.